

Module 2 Assignment 1

Ellen Bledsoe

2023-09-24

Assignment Details

Purpose

The goal of this assignment is to assess your ability to produce and interpret scatter plots in both base R and `ggplot2`.

Task

Write R code which produces the correct answers and text to correctly interpret the plots produced.

Criteria for Success

- Code is within the provided code chunks
- Code chunks run without errors
- Code produces the correct result
 - Code that produces the correct answer will receive full credit
 - Code attempts with logical direction will receive partial credit
- Written answers address the questions in sufficient detail

Due Date

Feb 24 before class

Assignment Questions

For this assignment, we are going to be making plots! We are going to use a data set containing data from a sample of our fish tanks. The data contains information about the tanks sampled and how many sick fish are contained in each tank.

First, let's prepare our data.

Load the `tidyverse` package into the work space.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.0     v stringr   1.5.2
## v ggplot2   4.0.0     v tibble    3.3.0
## v lubridate 1.9.4     v tidyv     1.3.1
## v purrr    1.1.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
```

```

## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
Read in the data using the read_csv() function. Name the data frame sick_fish.
sick_fish <- read_csv("../data/fish_sick_data.csv")

## Rows: 50 Columns: 7
## -- Column specification -----
## Delimiter: ","
## chr (1): species
## dbl (6): tank_id, avg_daily_temp, num_fish, day_length, tank_volume, num_sick
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

1. Use the `head()` function to take a look at the columns in the data frame.

How many rows does the data frame have? How many columns? (2 points)

Answer: 50 rows, 7 columns

```
head(sick_fish)
```

```

## # A tibble: 6 x 7
##   tank_id species avg_daily_temp num_fish day_length tank_volume num_sick
##   <dbl> <chr>        <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1     388 tilapia     24.3       93        10      399.       3
## 2     425 tilapia     24.6       98        11      400.       4
## 3     420 tilapia     23.0      103        9      399.       2
## 4     819 trout        14.1       85        11      401.      14
## 5     176 tilapia     23.3       98        10      400.       3
## 6     926 trout        13.8       79        12      400.      10

```

2. Take a look at the data frame. (2 points, 1 each)

- What does one row (observation) represent (e.g., an individual fish?, all fish of a certain species? all tanks of a certain species? etc.)?

Answer: one row represents one fish tank that was sampled

- Based on the column names, briefly describe what data is included in the data frame. (2 points)

Answer: information about each tank, the species of fish in it, and the environmental variables of the tank

3. Take a look at the columns that have the total number of fish in the tank and the number of sick fish per tank. Determine whether these two columns are numeric or categorical.

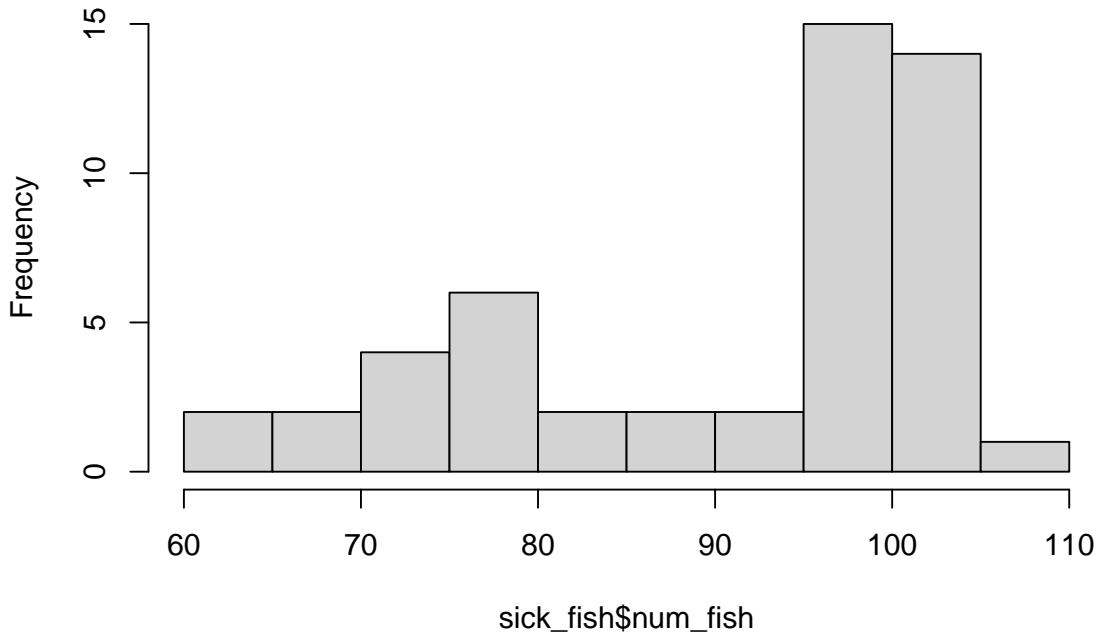
Answer: both `num_fish` and `num_sick` are numeric

Plotting with base R

4. Using *base R* (the code we learned in “Module2_1_DescriptiveStats”), create a histogram of the number of fish in each tank.

```
hist(sick_fish$num_fish)
```

Histogram of sick_fish\$num_fish



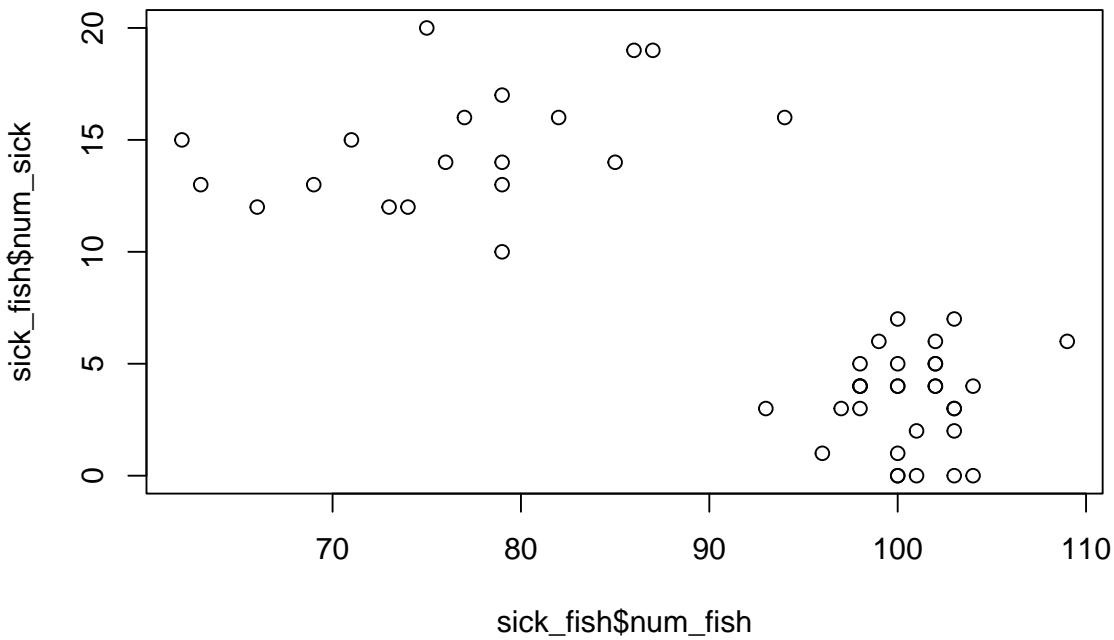
5. In 2-3 sentences, describe what the histogram tells you about the distribution of the number of sick fish. (2 points)

Note: they do not have to include exactly what I've typed below. As long as their answer is correct and is 2-3 sentences, you can give them full credit.

Answer: The range of the number of fish in each tank is between ~60-110. There are two main clusters: a small one between 70-80 and a much larger one between 95-105. The most frequent number of fish in tanks is between 95-100, with 15 tanks having that many fish.

6. Using *base R* (the code we learned in “Module2_1_DescriptiveStats”), create a scatter plot. Put the total number of fish on the x-axis and the number of sick fish on the y-axis.

```
plot(sick_fish$num_fish, sick_fish$num_sick)
```



- Interpret the scatter plot. What is the plot telling you? Is there a positive or negative relationship between the two variables. What does that mean? (2 points)

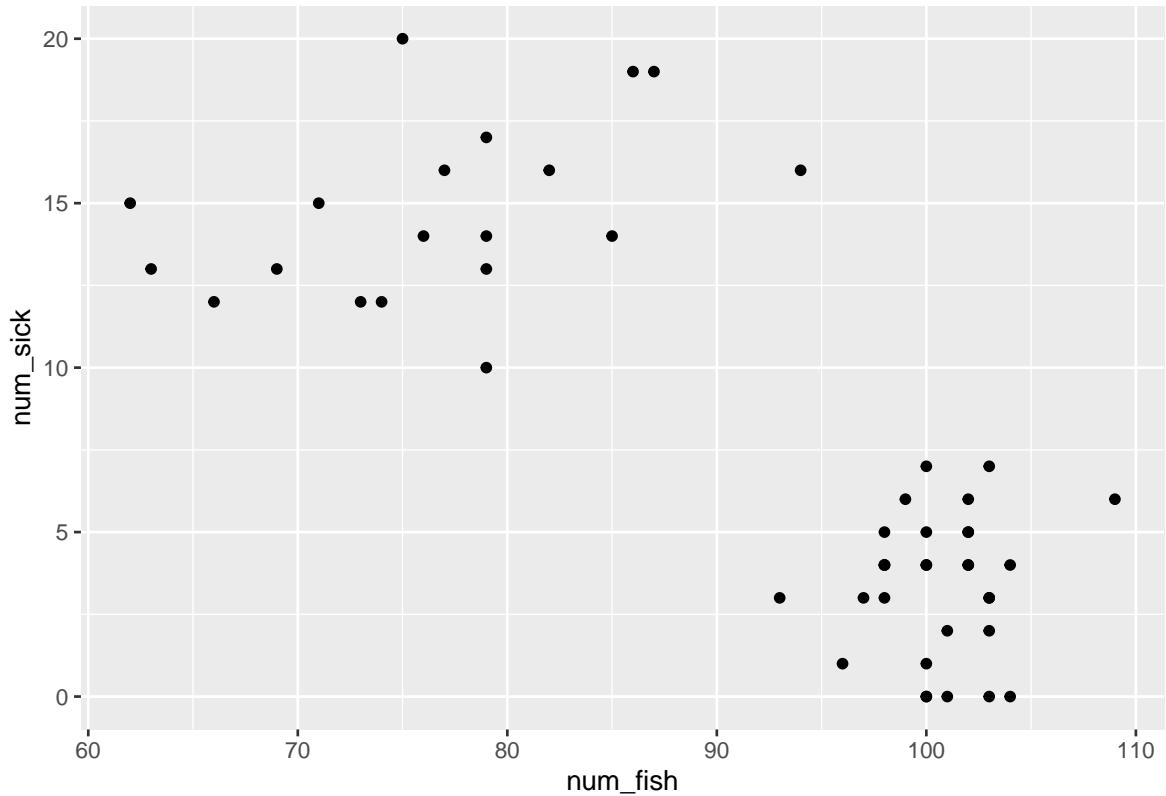
Answer: negative relationship, meaning that the more fish that are in the tank, the fewer sick fish there are (which is counter-intuitive to what me might predict, by the way).

Using `ggplot2`

Now, let's use what we've learned about `ggplot2` to recreate the same scatter plot. We will do this in 2 steps. In the first step, we will add the data, axes, and geom. In the second step, we will modify the plot to increase interpretability by renaming axes and adding a theme.

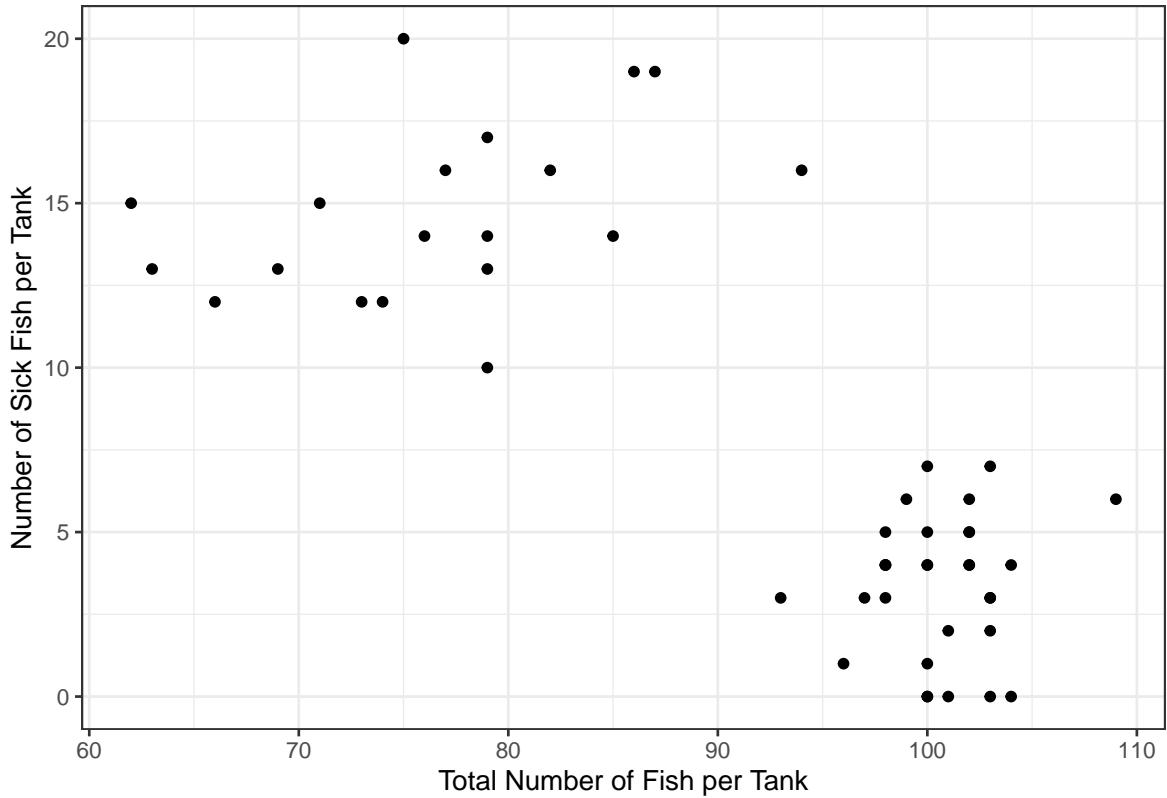
- Create the scatter plot using the proper `geom` function.

```
ggplot(sick_fish, aes(num_fish, num_sick)) +
  geom_point()
```



9. Now, add to the plot that we created above to make it clearer to understand. Add better axes labels and choose a theme. (2 point)

```
ggplot(sick_fish, aes(num_fish, num_sick)) +
  geom_point() +
  labs(x = "Total Number of Fish per Tank",
       y = "Number of Sick Fish per Tank") +
  theme_bw()
```



Looking at the scatter plots, there seem to be two distinct groups. Let's investigate this a bit further.

- For each fish species, calculate the average and the standard deviation for the number of sick fish. (2 points)

Hint: think back to Module 1 and which function we can use to “split, apply, combine”

```
sick_fish %>%
  group_by(species) %>%
  summarise(mean_sick = mean(num_sick),
            sd_sick = sd(num_sick))
```

```
## # A tibble: 2 x 3
##   species mean_sick sd_sick
##   <chr>     <dbl>   <dbl>
## 1 tilapia    3.39    2.11
## 2 trout      14.7    2.68
```

- Based on the summary data you've calculated above, do you think the two clumps of data correspond to the two species? Which species seems to be the one that is causing the most problems? (2 points)

Answer: seems like trout are driving the issues; seems reasonable that clumps could belong to species

- Do some brainstorming. Based on the plots we've made and the calculations we've done so far, are we *sure* that the two clumps of data we see in the data visualization do, in fact, correlate to the two different species? How could we confirm? How might we include that information in the plot? (2 points)

Answer: many answer options here. could take average of total num fish to see if that corresponds to what is on the plot; could change some characteristic of the plot (point shape, color, size) to represent fish species; could add labels (would get busy)